

System Threat Forecaster: Complete Technical Workflow

Machine Learning & Deep Learning Approaches

Technical Documentation

December 19, 2025

Agenda

- 1 Project Overview
- 2 Machine Learning Workflow
 - Environment Setup
 - Configuration
 - Utility Functions
 - Exploratory Data Analysis
 - Data Preprocessing
 - Feature Engineering & Selection
 - Model Training
 - ML Pipeline Execution
- 3 Deep Learning Workflow
 - PyTorch Setup
 - DL Configuration
 - PyTorch Dataset
 - Neural Network Architectures
 - Training Process
- 4 Comparison & Results
- 5 Web Application Integration
- 6 Conclusion

Project Overview

Objective

Predict potential malware infections in computer systems using machine learning and deep learning approaches

Dataset:

- 100,000 training samples
- 76 features (48 numeric, 28 categorical)
- Binary classification (malware vs clean)
- Balanced classes: 50.52% malware, 49.48% clean


Approaches:

- Traditional ML: 7 algorithms
- Deep Learning: 6 neural architectures
- Modular pipeline design
- Production-ready implementation

Key Features

- ① **Modular Architecture:** Independent components
- ② **Multiple Models:** Compare 7 ML + 6 DL algorithms
- ③ **Automated Pipeline:** End-to-end workflow
- ④ **Flexible Configuration:** Central CONFIG system
- ⑤ **Comprehensive Logging:** Track experiments
- ⑥ **Web App Integration:** Deployment-ready models

ML Workflow: Architecture Overview



ML: Environment Setup

```
1 import numpy as np, pandas as pd
2 from sklearn.preprocessing import StandardScaler, LabelEncoder
3 from sklearn.model_selection import train_test_split
4 from sklearn.metrics import classification_report, accuracy_score
5 from sklearn.ensemble import RandomForestClassifier
6 from lightgbm import LGBMClassifier
```

Key Libraries

Scikit-learn (ML algorithms), **LightGBM** (best performer), **Pandas/NumPy** (data), **Matplotlib** (visualization)

ML: Configuration System

```
1 CONFIG = {  
2     'data_path': {'train': './kaggle/input/train.csv',  
3                     'test': './kaggle/input/test.csv'},  
4     'run_eda': True,  
5     'models_to_train': {  
6         'decision_tree': True,  
7         'random_forest': True,  
8         'lightgbm': True, # Best performer  
9         'ada_boost': True  
0     },  
1     'random_state': 42  
2 }
```

Centralized Control

Single point to enable/disable pipeline components

Data Pipeline Control:

- File paths
- Missing value handling
- Feature engineering
- Feature selection
- Dimensionality reduction

Model Selection:

- 7 algorithms available
- Easy enable/disable
- Hyperparameter tuning
- Cross-validation folds
- Train-test split ratio

ML: Utility Functions

```
1 def load_data(path):
2     return pd.read_csv(path)
3
4 def save_model(model, model_name):
5     timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")
6     filename = f"models/{model_name}_{timestamp}.joblib"
7     joblib.dump(model, filename)
8
9 def log_result(model_name, metrics, hyperparams=None):
10    log_data = { 'model_name': model_name,
11                 'accuracy': metrics['accuracy'],
12                 'timestamp': datetime.now()}
13    df.to_csv('results/model_comparison.csv', mode='a')
```

- ① **Data Management:** Consistent loading, error handling
- ② **Model Persistence:** Timestamped files for versioning
- ③ **Results Logging:** Performance tracking and model selection

Purpose

Understand dataset characteristics to guide preprocessing and model selection

Dataset Analysis:

- Shape: $100,000 \times 76$
- Data types identification
- Missing value patterns
- Statistical summaries

Feature Analysis:

- Distribution plots
- Correlation heatmap
- Target variable balance
- Top correlated features

ML: Key EDA Findings

- **Features:** 48 numeric, 28 categorical
- **Target:** Balanced (50.52% vs 49.48%)
- **Missing Values:** Several columns affected
- **Top Correlations:** AntivirusConfigID (0.118), TotalPhysicalRAMMB (0.066)
- **Insight:** Moderate correlations → ensemble models needed

ML: Preprocessing Pipeline

- ① **Feature Separation:** Split X, y and identify column types
- ② **Missing Values:** Mean (numeric), mode (categorical)
- ③ **Encoding:** LabelEncoder for categorical features
- ④ **Scaling:** StandardScaler: $z = \frac{x-\mu}{\sigma}$

ML: Preprocessing Code

```
def preprocess_data(data, is_training=True):
    X = data.drop('target', axis=1)
    y = data['target']

    numeric_cols = X.select_dtypes(include=['int64']).columns
    categorical_cols = X.select_dtypes(include=['object']).columns

    X[numeric_cols] = SimpleImputer(strategy='mean').fit_transform(X[
        numeric_cols])
    X[categorical_cols] = SimpleImputer(strategy='most_frequent').
        fit_transform(X[categorical_cols])

    X = LabelEncoder().fit_transform(X)
    X_scaled = StandardScaler().fit_transform(X)

    return X_scaled, y
```

ML: Optional Feature Processing

Feature Engineering (Disabled)

Creates interaction features (e.g., AntivirusConfigID \times TotalPhysicalRAMMB)

Feature Selection (Disabled)

SelectKBest with ANOVA F-test keeps top 30 features

Dimensionality Reduction (Disabled)

PCA with 95% variance \rightarrow ~ 40 components

ML: Seven Algorithms Implemented

Model	Accuracy	Type
LightGBM	63.0%	Gradient Boosting
Random Forest	62.4%	Ensemble (Bagging)
AdaBoost	62.0%	Ensemble (Boosting)
Decision Tree	~60%	Single Tree
Logistic Regression	~60%	Linear Model
Naive Bayes	Lower	Probabilistic
SGD	Lower	Online Learning

Best Model

LightGBM achieves 63.0% accuracy due to:

- Leaf-wise tree growth (vs level-wise)
- Handles large datasets efficiently
- Built-in categorical feature support

ML: Training Process

```
1 def train_model(model_name, X_train, y_train, X_val, y_val):
2     params = get_default_model_params(model_name)
3
4     model = get_model(model_name, params)
5     model.fit(X_train, y_train)
6
7     val_pred = model.predict(X_val)
8     metrics = {
9         'val_accuracy': accuracy_score(y_val, val_pred),
10        'precision': precision_score(y_val, val_pred),
11        'f1_score': f1_score(y_val, val_pred)
12    }
13
14    save_model(model, model_name)
15    return model, metrics
```

ML: Evaluation Metrics

Accuracy:

$$\frac{\text{Correct Predictions}}{\text{Total Predictions}}$$

Recall:

$$\frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

Precision:

$$\frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

F1-Score:

$$2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

Confusion Matrix

Visualizes True Positives, False Positives, True Negatives, False Negatives


ML: Complete Pipeline Flow

- ① Load train.csv and test.csv
- ② Run EDA and preprocess (handle missing, encode, scale)
- ③ Split train (80%) / validation (20%)
- ④ Optional: Feature engineering, selection, PCA
- ⑤ Train all selected models (7 algorithms)
- ⑥ Compare performance and select best
- ⑦ Generate predictions and save models
- ⑧ Export metrics to results/model_performance.json

ML: Output Files

File	Purpose
saved_models/ml_models.pkl	All trained models
saved_models/preprocessors.pkl	Scalers, encoders
results/model_comparison.csv	Performance log
results/model_performance.json	Structured metrics
submission.csv	Kaggle submission
models/lightgbm_*.joblib	Versioned models

DL Workflow: Architecture Overview



DL: Environment Setup

```
1 import torch, torch.nn as nn
2 from torch.utils.data import Dataset, DataLoader
3
4 torch.manual_seed(42) # Reproducibility
5
6 # Device detection (GPU acceleration)
7 if torch.cuda.is_available():
8     device = torch.device('cuda') # NVIDIA
9 elif torch.backends.mps.is_available():
10    device = torch.device('mps') # Apple Silicon
11 else:
12     device = torch.device('cpu') # CPU
```

GPU Acceleration

CUDA (10-100×), MPS (5-15×), CPU (1×)

DL: Configuration

```
1 CONFIG = {  
2     'batch_size': 512,  
3     'epochs': 100,  
4     'learning_rate': 0.001,  
5     'hidden_dims': [256, 128, 64, 32],  
6     'dropout_rate': 0.3,  
7     'use_batch_norm': True,  
8     'optimizer': 'adamw',  
9     'use_mixed_precision': torch.cuda.is_available()  
0 }
```

DL: Configuration Comparison

Setting	ML	DL
Training	One-shot .fit()	Iterative epochs
Batch Size	N/A	512
Learning Rate	N/A	0.001
Regularization	N/A	Dropout (0.3)
Data Format	DataFrame	Tensors
Feature Eng.	Manual	Automatic
Computation	CPU	GPU-accelerated
Training Time	Minutes	Hours

DL: Custom Dataset Class

```
1 class TabularDataset(Dataset):
2     def __init__(self, X, y=None):
3         self.X = torch.FloatTensor(X)
4         self.y = torch.LongTensor(y) if y else None
5
6     def __len__(self): return len(self.X)
7     def __getitem__(self, idx):
8         return (self.X[idx], self.y[idx]) if self.y else self.X[idx]
9
10 train_loader = DataLoader(train_dataset, batch_size=512,
11                           shuffle=True, num_workers=4)
```

- ① **Batching:** Groups 512 samples, memory efficient
- ② **Shuffling:** Prevents order patterns, improves generalization
- ③ **Parallel Loading:** 4 workers eliminate bottleneck
- ④ **Class Imbalance:** Weighted loss [1.98, 2.02]

DL: Six Architectures

Model	Accuracy	Time	Params
Simple MLP	60-62%	5 min	50K
Deep MLP	62-64%	10 min	100K
Residual Net	63-65%	15 min	200K
Attention Net	64-66%	20 min	300K
Wide & Deep	63-65%	12 min	150K
FT-Transformer	65-68%	30 min	500K

Best Architecture

FT-Transformer: 4-5% improvement over traditional ML

Architecture 1: Simple MLP

Structure:

- Input (76 features)
- Linear($76 \rightarrow 256$) + ReLU + Dropout
- Linear($256 \rightarrow 128$) + ReLU + Dropout
- Linear($128 \rightarrow 64$) + ReLU + Dropout
- Output($64 \rightarrow 2$)

Components:

- **Linear:** $y = Wx + b$
- **ReLU:** $\max(0, x)$
- **Dropout:** Zeros 30% neurons

Use Case:

- Baseline model
- Fast training
- 60-62% accuracy

Architecture 2: Deep MLP with Batch Norm

```
1 class DeepMLP(nn.Module):
2     def __init__(self, input_dim, hidden_dims=[256, 128, 64]):
3         super().__init__()
4         layers = []
5         for hidden_dim in hidden_dims:
6             layers.append(nn.Linear(prev_dim, hidden_dim))
7             layers.append(nn.BatchNorm1d(hidden_dim))
8             layers.append(nn.ReLU())
9             layers.append(nn.Dropout(0.3))
10        self.network = nn.Sequential(*layers)
```


Batch Normalization

$$x_{norm} = \frac{x - \mu_{batch}}{\sqrt{\sigma_{batch}^2 + \epsilon}}, \quad output = \gamma \cdot x_{norm} + \beta$$

Faster training, stability (62-64% accuracy)

Architecture 3: Residual Network

Skip Connections:



$$out = Block(x) + x$$

Benefits:

- Solves vanishing gradients
- Deeper networks (100+ layers)
- 63-65% accuracy

Architecture 4: Attention Network

Self-Attention: $\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) \cdot V$

Mechanism:

- Each feature attends to others
- Learns interactions automatically

Benefits:

- Complex relationships
- Interpretable weights
- 64-66% accuracy

Innovation

Transforms NLP architecture for tabular data

Architecture 5: Wide & Deep

Simplified Architecture:

Input (76 features) splits into:

- Wide path: Direct Linear($76 \rightarrow 2$)
- Deep path: MLP $256 \rightarrow 128 \rightarrow 64 \rightarrow 2$

Both outputs are added together for final prediction.

Components:

- **Wide:** Memorization (frequent patterns)
- **Deep:** Generalization (novel combinations)

Use Case:

- Google Play recommendations
- 63-65% accuracy

Architecture 6: FT-Transformer (Best)

Feature Tokenizer Transformer - SOTA for tabular data

- ① **Tokenization:** Each feature → 192-dim vector
- ② **CLS Token:** Prepend classification token (like BERT)
- ③ **Positional Embeddings:** Learnable position encoding
- ④ **Transformer:** 8-head self-attention, 3 blocks
- ⑤ **Classification:** CLS token → Linear($192 \rightarrow 2$)

Performance

65-68% accuracy - 4-5% improvement over ML

DL: Training Loop

- ① **Forward:** $\hat{y} = f_{\theta}(x)$ (network prediction)
- ② **Loss:** $L = -\log(P(\text{correct class}))$ (CrossEntropy)
- ③ **Backward:** $\frac{\partial L}{\partial \theta}$ (compute gradients)
- ④ **Update:** $\theta_{t+1} = \theta_t - \alpha \cdot \nabla_{\theta} L$ (AdamW)

Iterative Learning

Repeat for 100 epochs

DL: Training Code

```
1 def train_epoch(model, train_loader, criterion, optimizer):
2     model.train()
3
4     for data, target in train_loader:
5         output = model(data)
6         loss = criterion(output, target)
7
8         optimizer.zero_grad()
9         loss.backward()
10        optimizer.step()
11
12    return avg_loss, accuracy
13
14 for epoch in range(100):
15     train_loss = train_epoch(...)
16     val_loss = validate(...)
17     if early_stopping(val_loss): break
```

- ① **Early Stopping:** Stop if no improvement for 15 epochs
- ② **LR Scheduling:** Reduce when stuck ($0.001 \rightarrow 0.0005$)
- ③ **Mixed Precision:** FP16 for 2-3 \times speedup (NVIDIA)
- ④ **Checkpointing:** Save best model by validation accuracy

DL: Complete Pipeline Flow

- ① Load and preprocess → NumPy arrays
- ② Create DataLoaders (batch size=512)
- ③ Split train (80%) / validation (20%)
- ④ Train 6 architectures (MLP, ResNet, Attention, FT-Transformer)
- ⑤ Evaluate and select best model (FT-Transformer ~67%)
- ⑥ Generate predictions and save models
- ⑦ Export to submission_dl.csv

ML vs DL: Comprehensive Comparison

Aspect	Machine Learning	Deep Learning
Best Model	LightGBM (63.0%)	FT-Transformer (65-68%)
Training	One-shot .fit()	Iterative epochs
Data Format	DataFrame	Tensors
Feature Eng.	Manual	Automatic
Architecture	Fixed (trees)	Flexible (networks)
Computation	CPU-optimized	GPU-accelerated
Training Time	5-15 minutes	30-60 minutes
Interpretability	High	Low (black box)
Hyperparams	Few (~10)	Many (~50)
Memory	Low	High (GPU)

Performance Summary

Category	Model	Accuracy
Traditional ML	LightGBM	63.0%
	Random Forest	62.4%
	AdaBoost	62.0%
Deep Learning	FT-Transformer	65-68%
	Attention Net	64-66%
	Residual Net	63-65%

Key Finding

Deep Learning achieves **4-5% improvement** over traditional ML

When to Use ML vs DL

Use Traditional ML When:

- Small dataset ($\leq 10K$ samples)
- Need interpretability
- Limited compute resources
- Quick prototyping
- Tabular data with clear features
- Need fast inference

Use Deep Learning When:

- Large dataset ($\geq 50K$ samples)
- Complex patterns
- GPU available
- Maximum performance needed
- Many feature interactions
- Time for experimentation

Our Project

With 100K samples and GPU access, DL provides the best performance

Deployment Pipeline

- ① **Model Saving:** ML (7 models), DL (6 architectures), preprocessors
- ② **Metadata:** Best model, accuracy, hyperparameters
- ③ **Web App:** Load preprocessors and best model (FT-Transformer)
- ④ **Inference:** User input → Preprocess → Model → Prediction

Web App Integration Code

```
1 import joblib, torch
2
3 preprocessors = joblib.load('saved_models/preprocessors.pkl')
4 checkpoint = torch.load('saved_models/dl_models.pth')
5 model = checkpoint['best_model']
6 model.eval()
7
8 def predict(input_features):
9     X = preprocessors['scaler'].transform([input_features])
10
11     with torch.no_grad():
12         output = model(torch.FloatTensor(X))
13         pred = output.argmax(dim=1).item()
14
15     return "Malware" if pred == 1 else "Clean"
```

Key Achievements

- ① **Implementation:** 7 ML + 6 DL algorithms, production-ready
- ② **Performance:** ML (63.0%), DL (65-68%), 4-5% improvement
- ③ **Technical:** GPU acceleration, mixed precision, Transformers
- ④ **Deployment:** Model versioning, web app, comprehensive logging

Technical Highlights

Machine Learning

- Modular configuration system
- Multiple algorithm comparison
- Comprehensive preprocessing pipeline
- Feature engineering capabilities

Deep Learning

- 6 state-of-the-art architectures
- Transformer-based approach for tabular data
- Advanced training techniques (early stopping, LR scheduling)
- GPU acceleration support

Future Improvements

- ① **Ensemble:** Combine ML + DL predictions, stacking
- ② **Optimization:** Optuna, Ray Tune, NAS
- ③ **Features:** Cross-validation, advanced imputation
- ④ **Deployment:** ONNX export, quantization, A/B testing

Thank You

Questions?

Repository Structure:

- system-threat-forecaster-ml.py - ML Pipeline
- system-threat-forecaster-dl.py - DL Pipeline
- saved_models/ - Trained models
- results/ - Performance metrics
- models/ - Model checkpoints

Best Performance: FT-Transformer at 65-68% accuracy

Appendix: Hyperparameters - LightGBM

Parameter	Value
n_estimators	200
learning_rate	0.1
max_depth	5
subsample	0.6
colsample_bytree	0.8
min_child_samples	20
reg_alpha	0.1
reg_lambda	0.1

Appendix: Hyperparameters - FT-Transformer

Parameter	Value
d_token	192
n_blocks	3
attention_heads	8
attention_dropout	0.2
ffn_dropout	0.1
residual_dropout	0.0
learning_rate	0.001
weight_decay	1e-5
batch_size	512
epochs	100

Appendix: Dataset Features

Numeric Features (48):

- TotalPhysicalRAMMB
- AvailablePhysicalRAMMB
- ProcessorCount
- SystemVolumeTotalCapacity
- SystemDriveFreeSpace
- ...

Categorical Features (28):

- AntivirusConfigID
- FirewallConfigID
- OSVersion
- CountryIdentifier
- LocaleEnglishName
- ...

Target Variable: Binary (0=Clean, 1=Malware)